

CSCI 4830 / 5722

Computer Vision



University of Colorado **Boulder**

CSCI 4830 / 5722

Computer Vision



Dr. Ioana Fleming
Spring 2019
Lecture 9



University of Colorado **Boulder**

Reminders

Submissions:

- Homework 2: due Wed 2/13 at 11 pm
- Homework 3: later this week

Readings:

- Szeliski:
 - chapter 3 (filters, changing resolution, Laplacian pyramids, warping)
 - chapter 4.1 (points) and 4.2 (edge detection)
- P&F Ch. 4,5



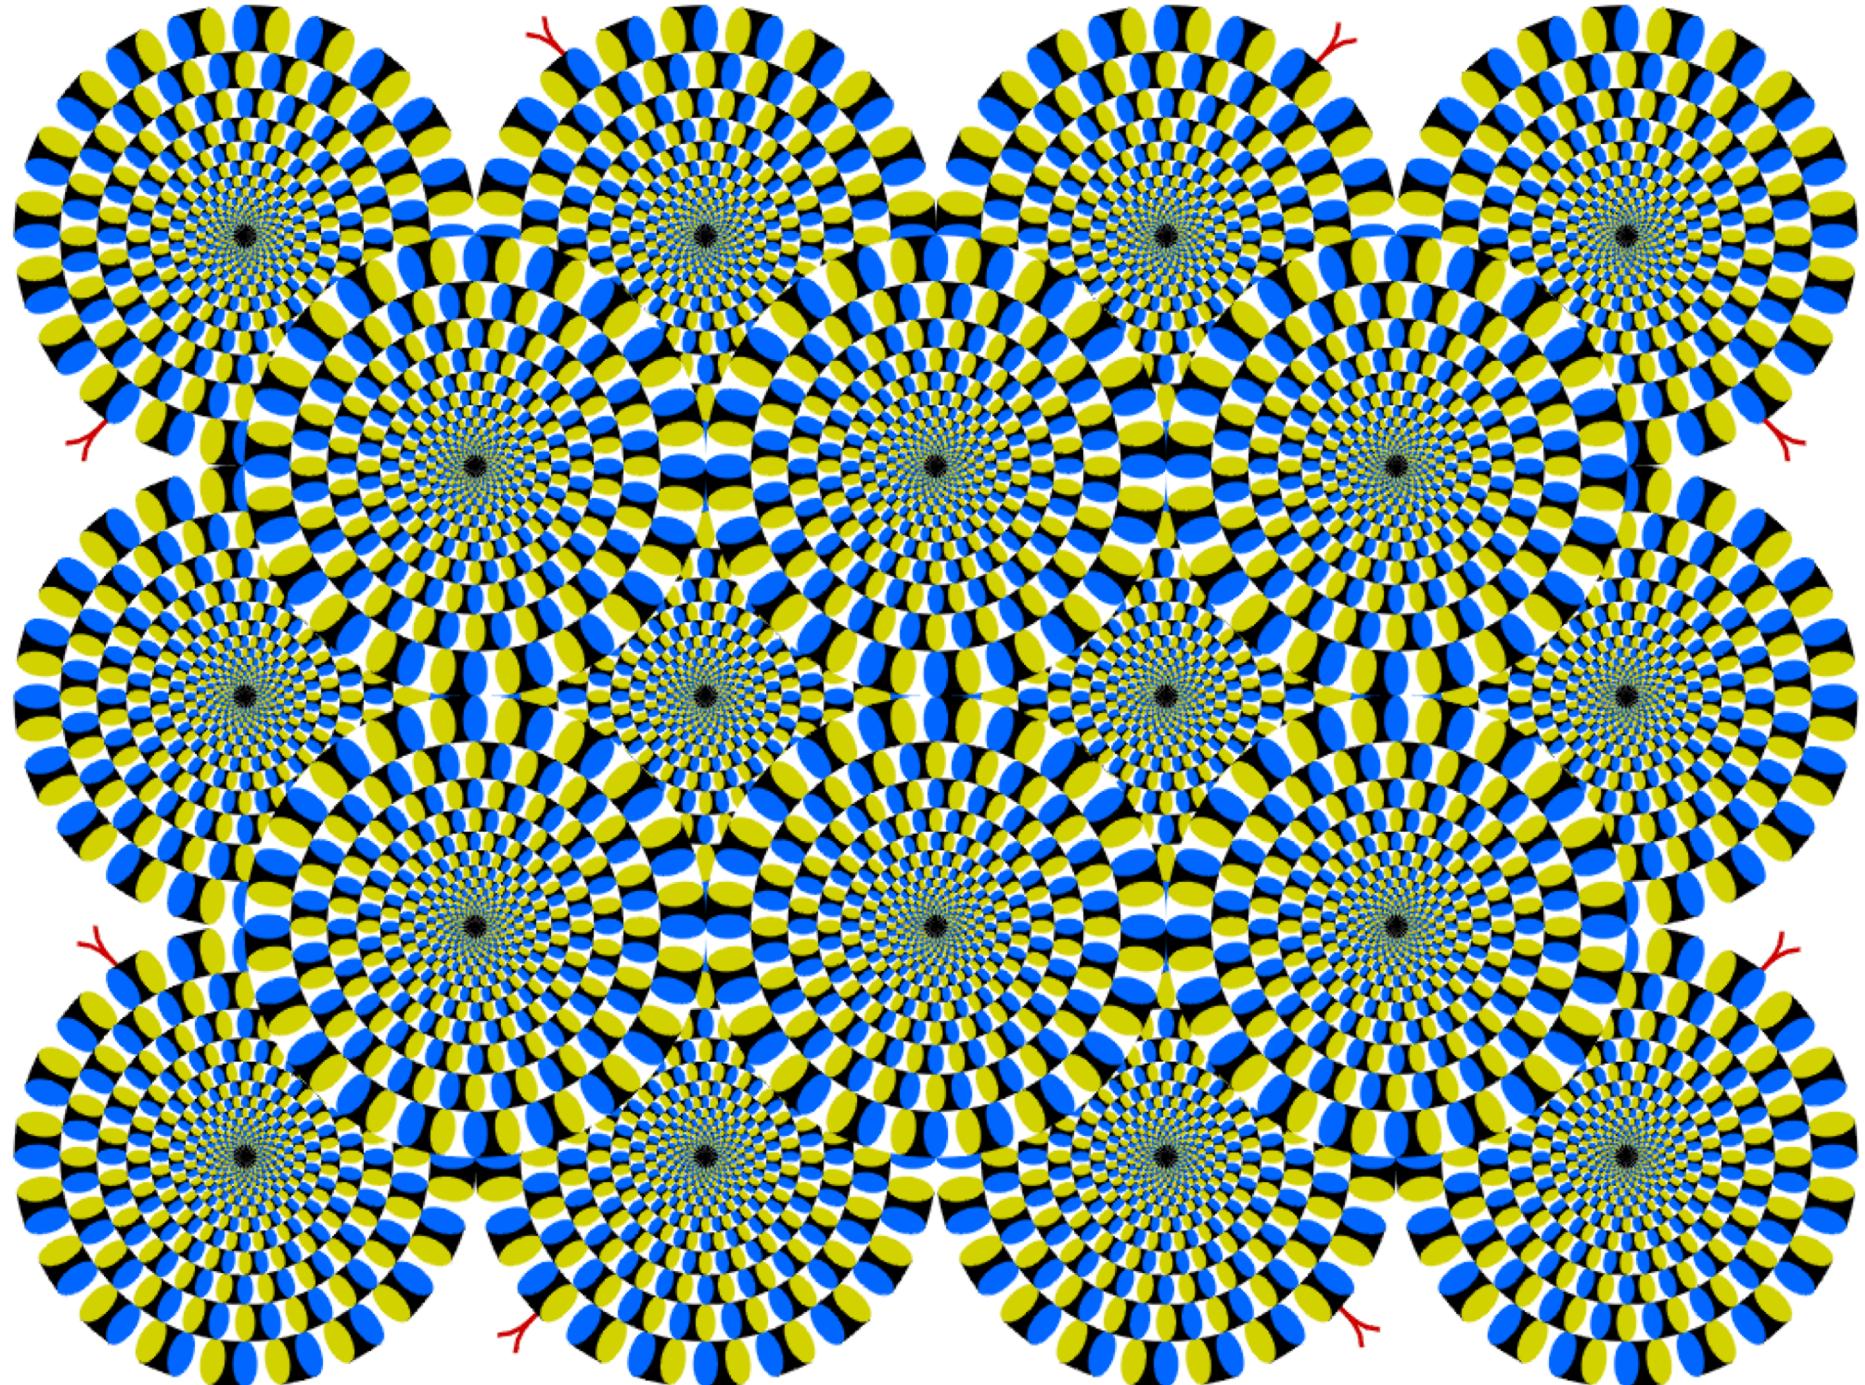
University of Colorado **Boulder**

Today

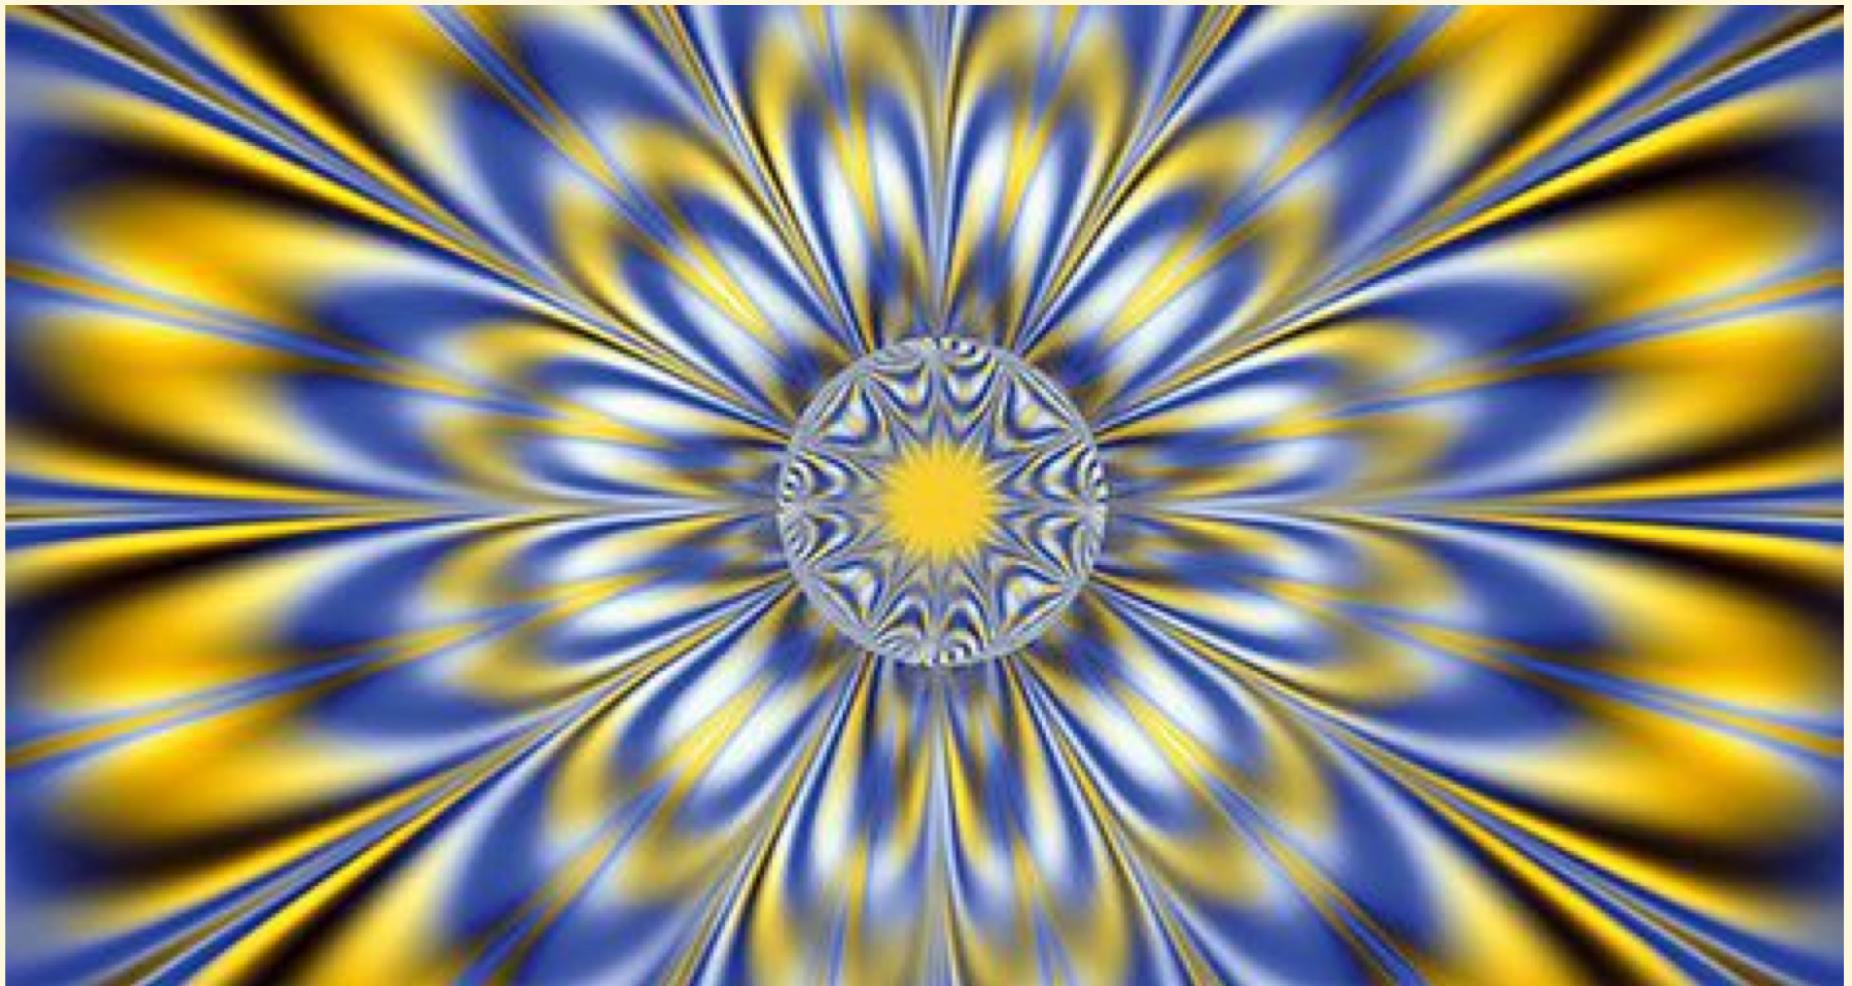
- Intensity Surfaces
- Gradients
- Image filters as linear operators



University of Colorado **Boulder**



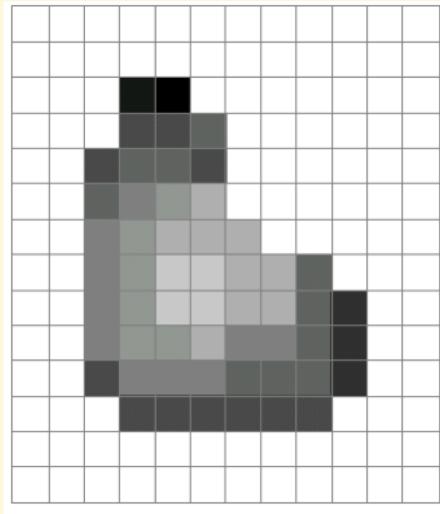
Moiré effect



University of Colorado **Boulder**

What is an image?

- A grid (matrix) of intensity values



=

255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	20	0	255	255	255	255	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255	255	255	255	255
255	255	127	145	145	175	127	127	95	47	255	255	255	255	255	255
255	255	74	127	127	127	95	95	95	47	255	255	255	255	255	255
255	255	255	74	74	74	74	74	74	74	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255

(common to use one byte per value: 0 = black, 255 = white)



University of Colorado **Boulder**

Visualizing Images

Our goal: we want to visualize images at a level high enough to retain human insight, but low enough to allow us to readily translate our insights into mathematical notation and, ultimately, computer algorithms that operate on arrays of numbers.



What is an image?

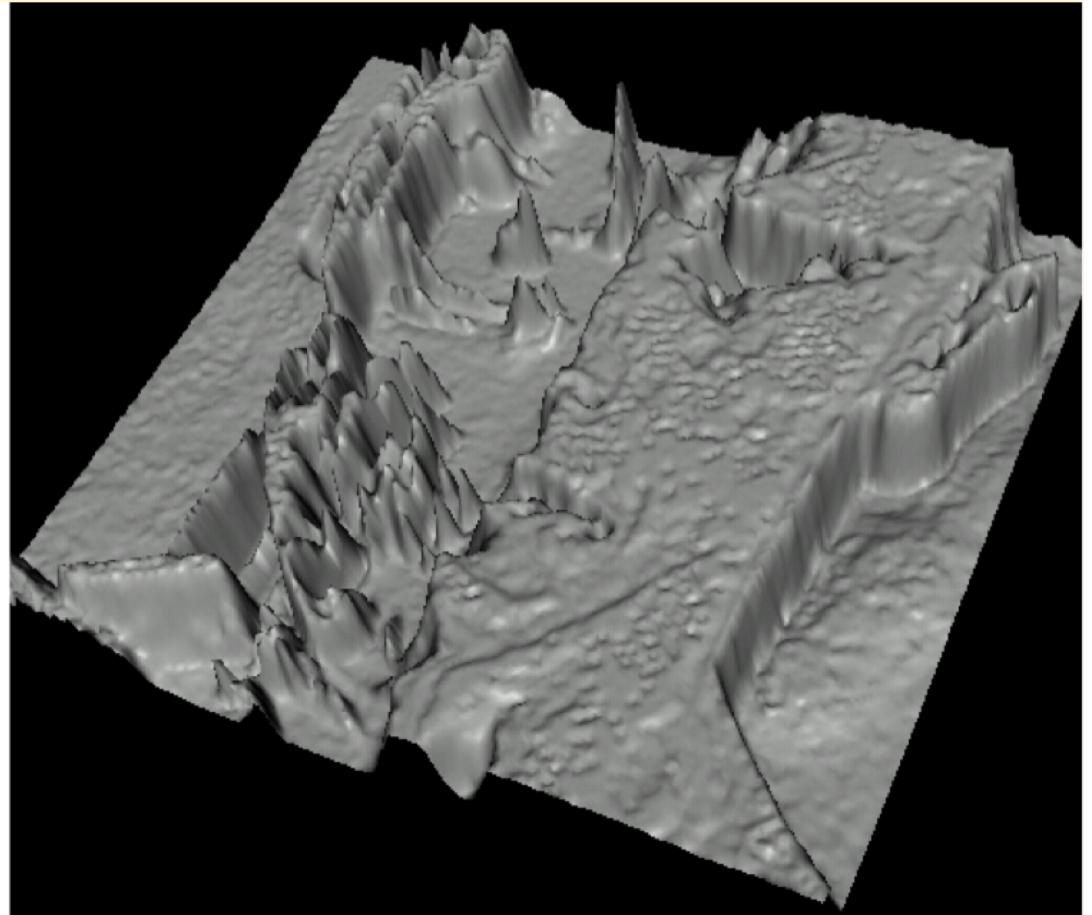
- We can think of a (grayscale) image as a **function**, f , from \mathbb{R}^2 to \mathbb{R} :
 - $f(x,y)$ gives the **intensity** at position (x,y)

A **digital** image is a discrete
(sampled, quantized) version of
this function



University of Colorado **Boulder**

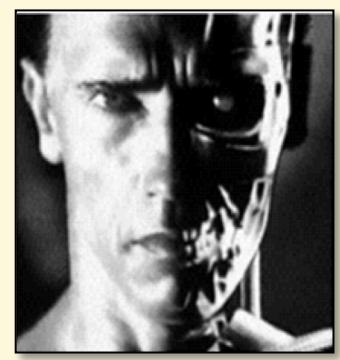
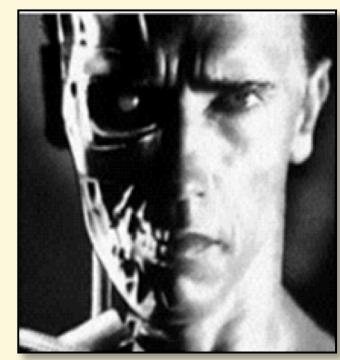
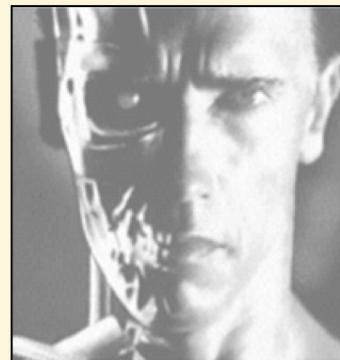
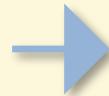
What is an image?



University of Colorado **Boulder**

Image transformations

- As with any function, we can apply operators to an image



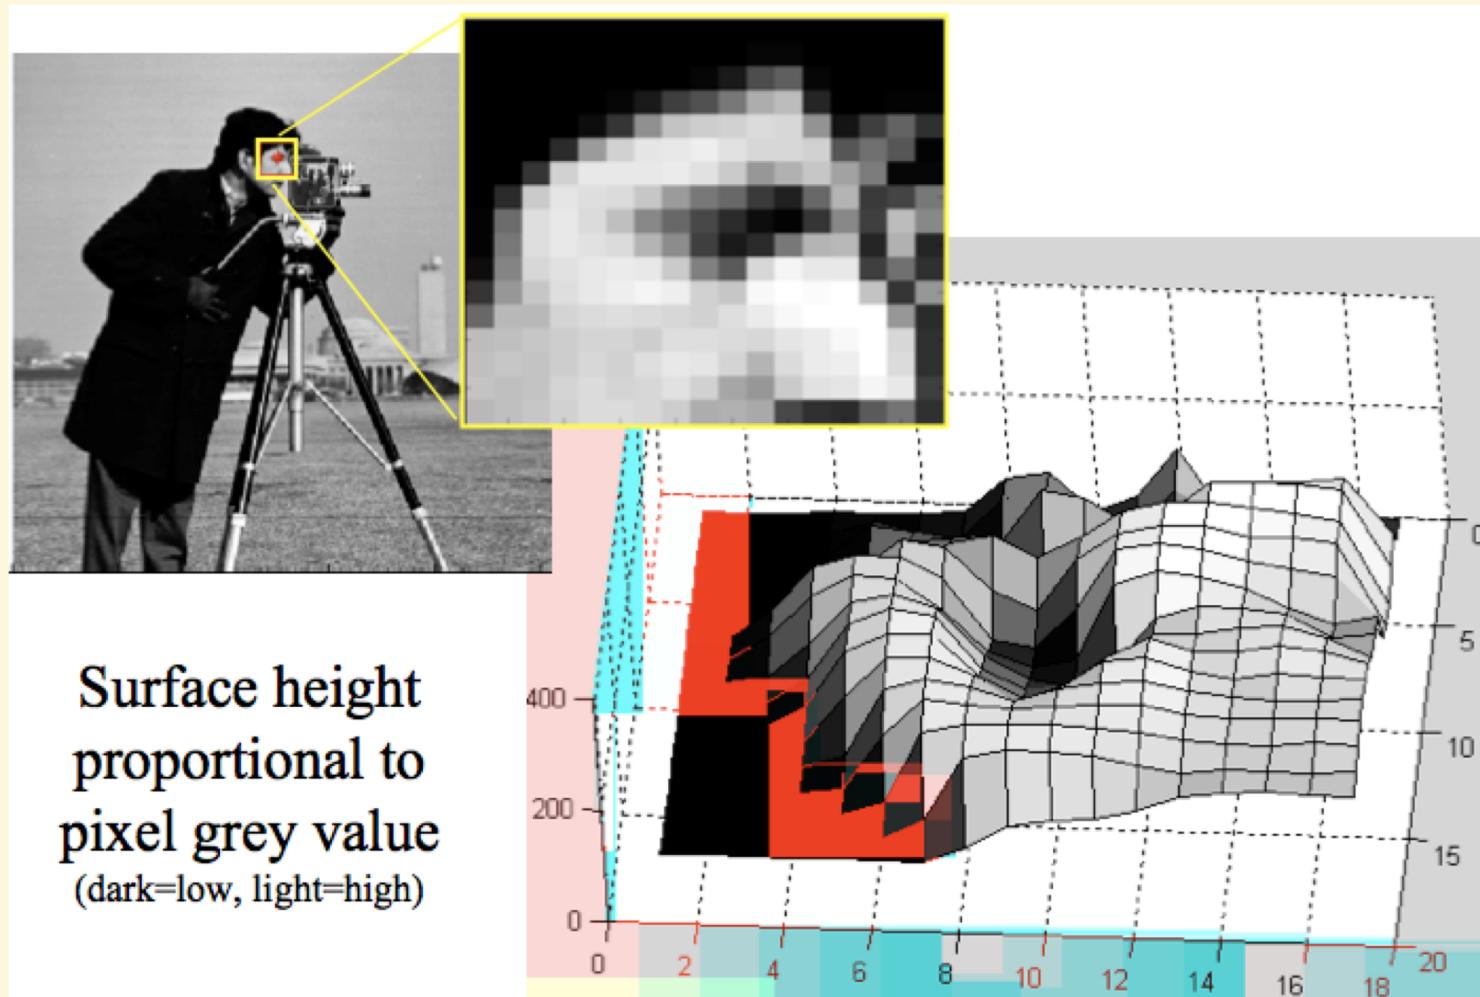
$$g(x,y) = f(x,y) + 20$$

$$g(x,y) = f(-x,y)$$



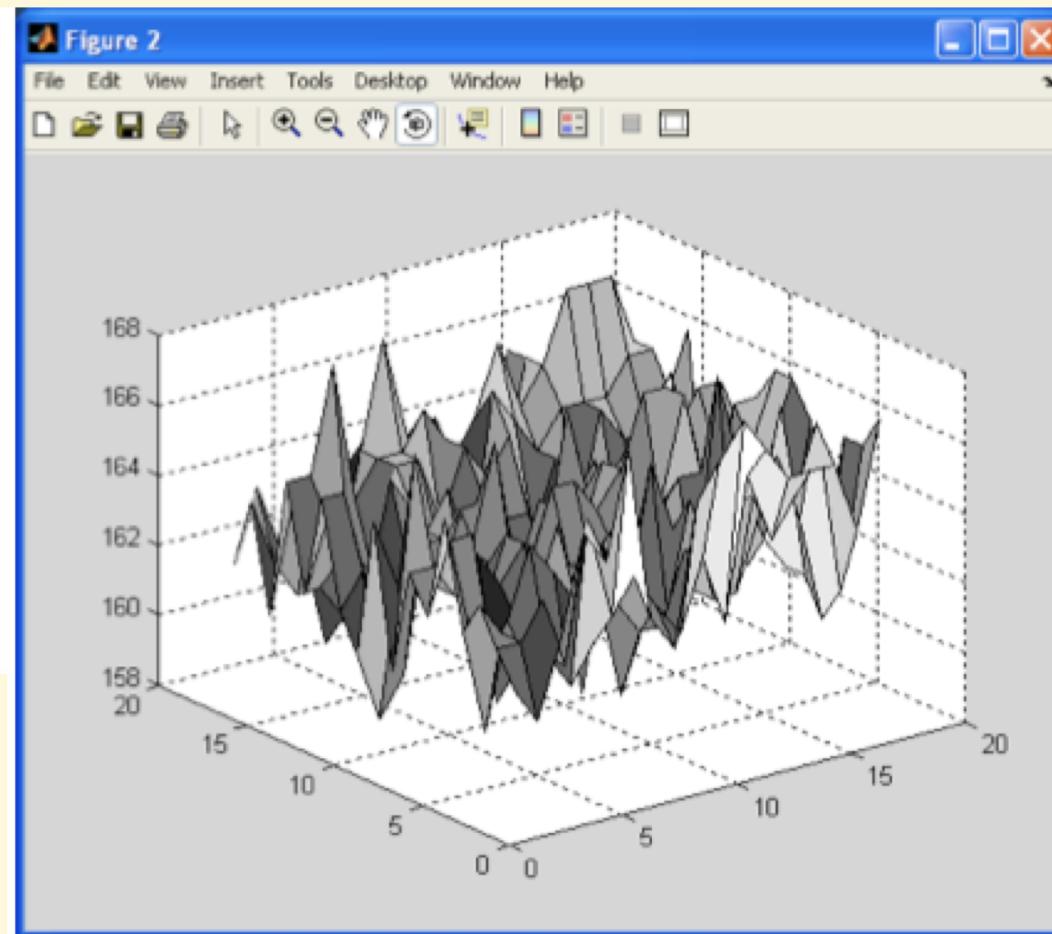
University of Colorado **Boulder**

Images as surfaces



University of Colorado **Boulder**

Images as surfaces

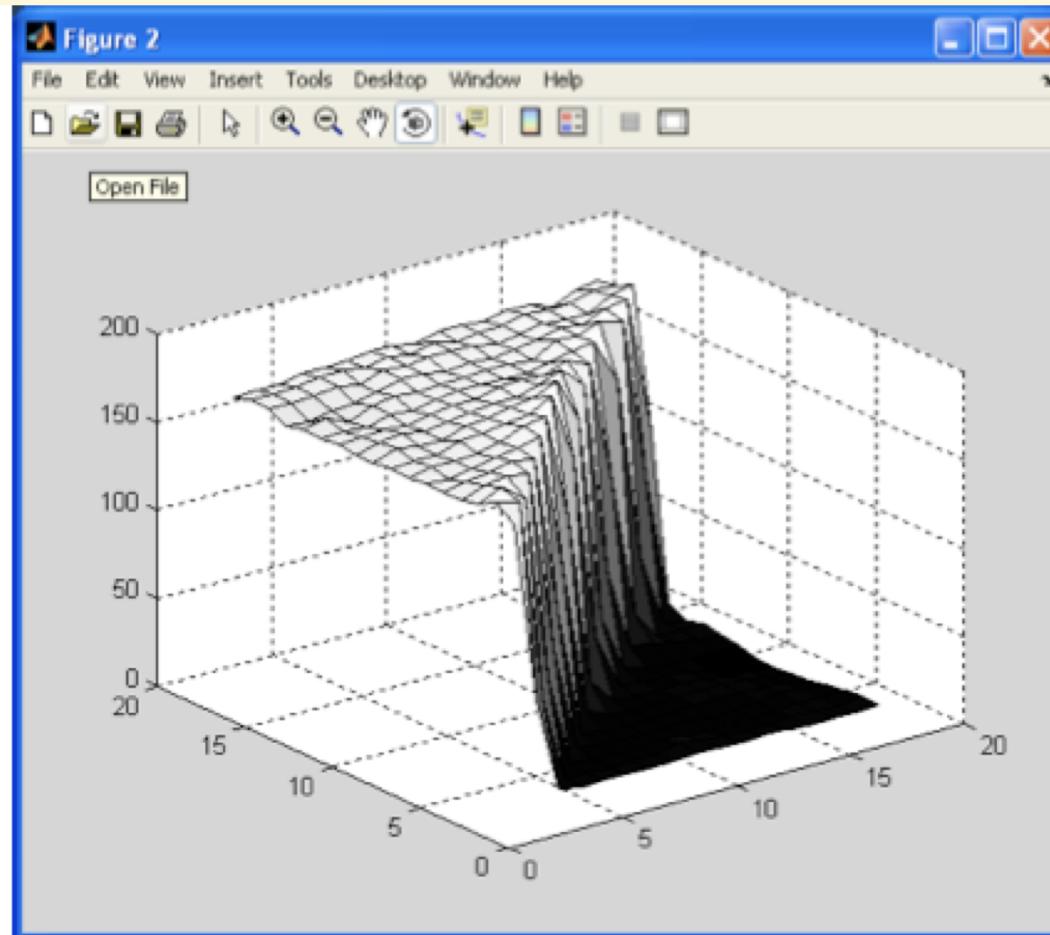


Mean = 164 Std = 1.8



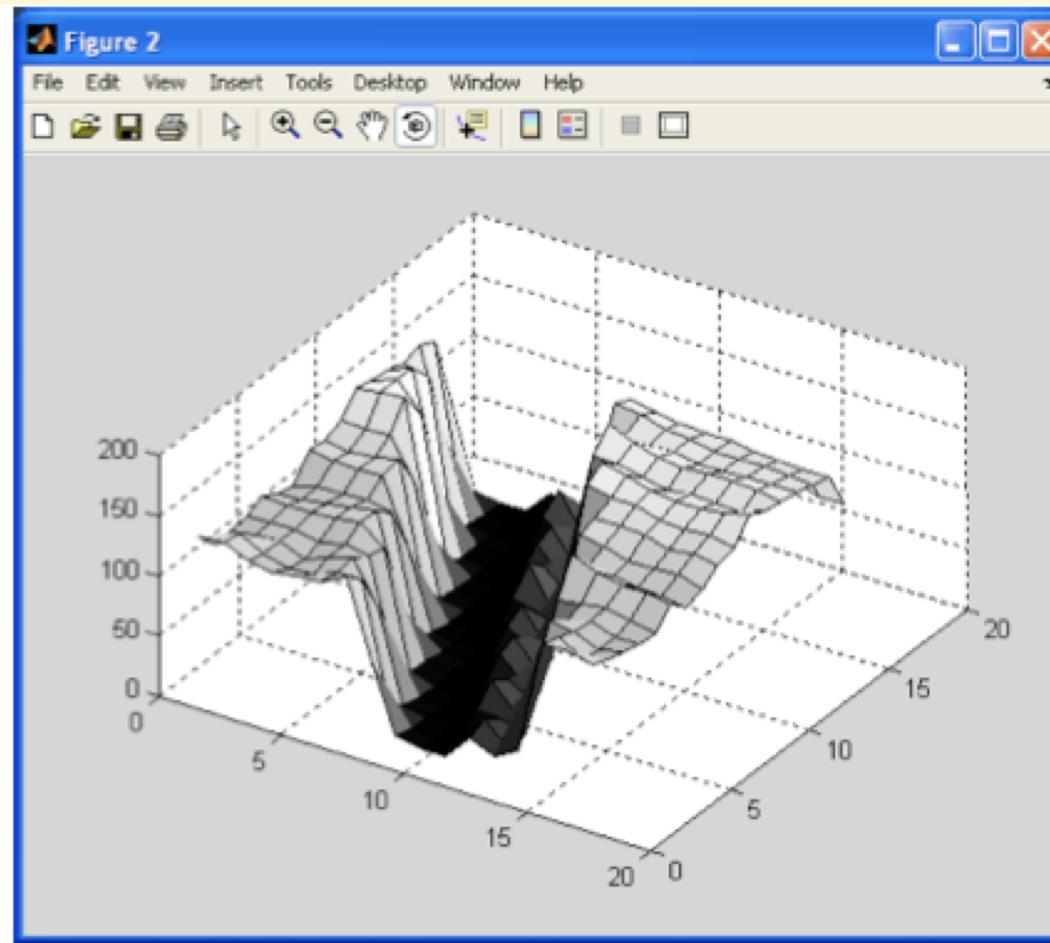
University of Colorado **Boulder**

Images as surfaces



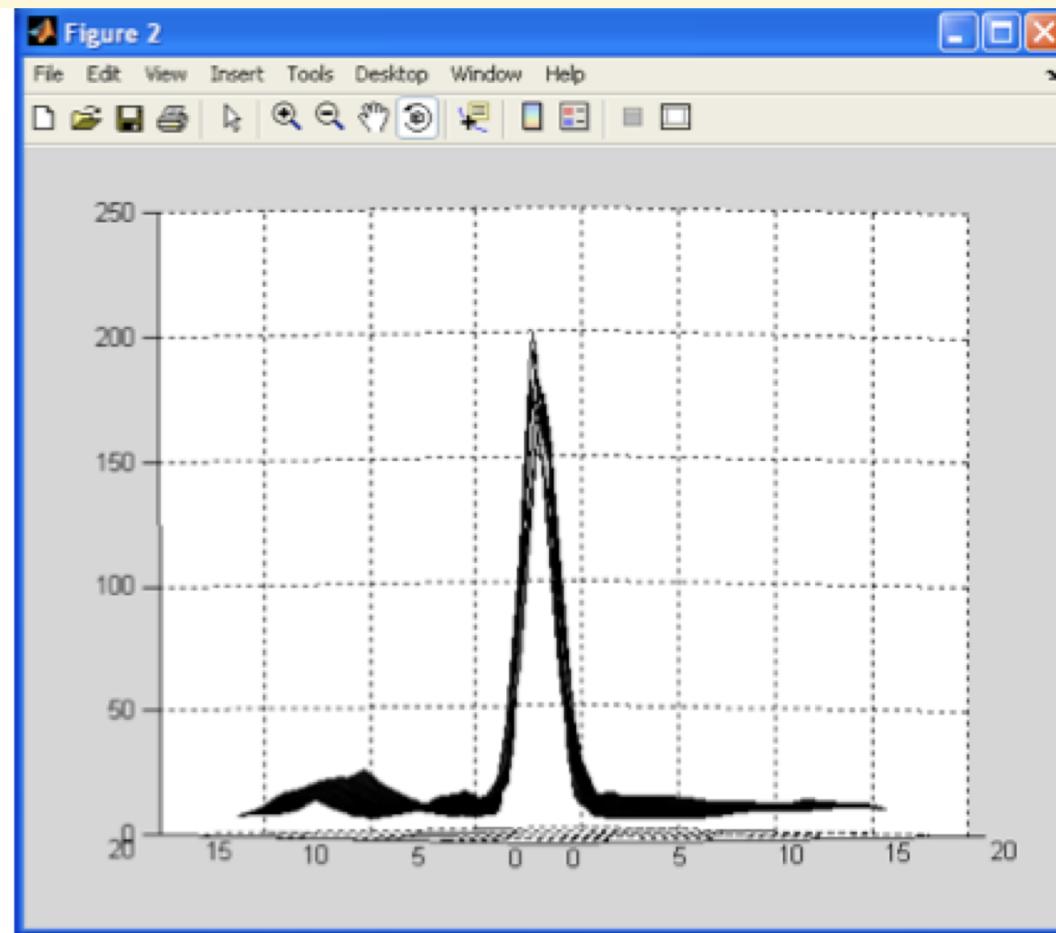
University of Colorado **Boulder**

Images as surfaces



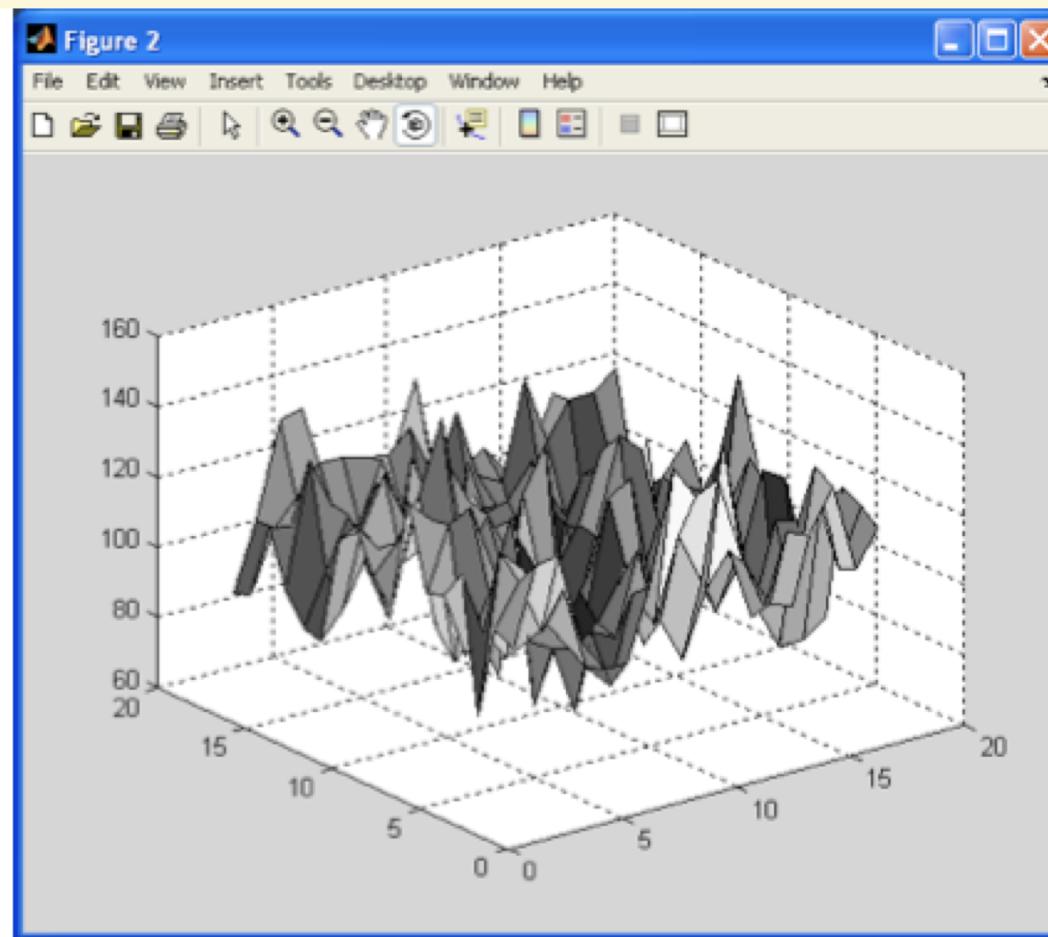
University of Colorado **Boulder**

Images as surfaces



University of Colorado **Boulder**

Images as surfaces



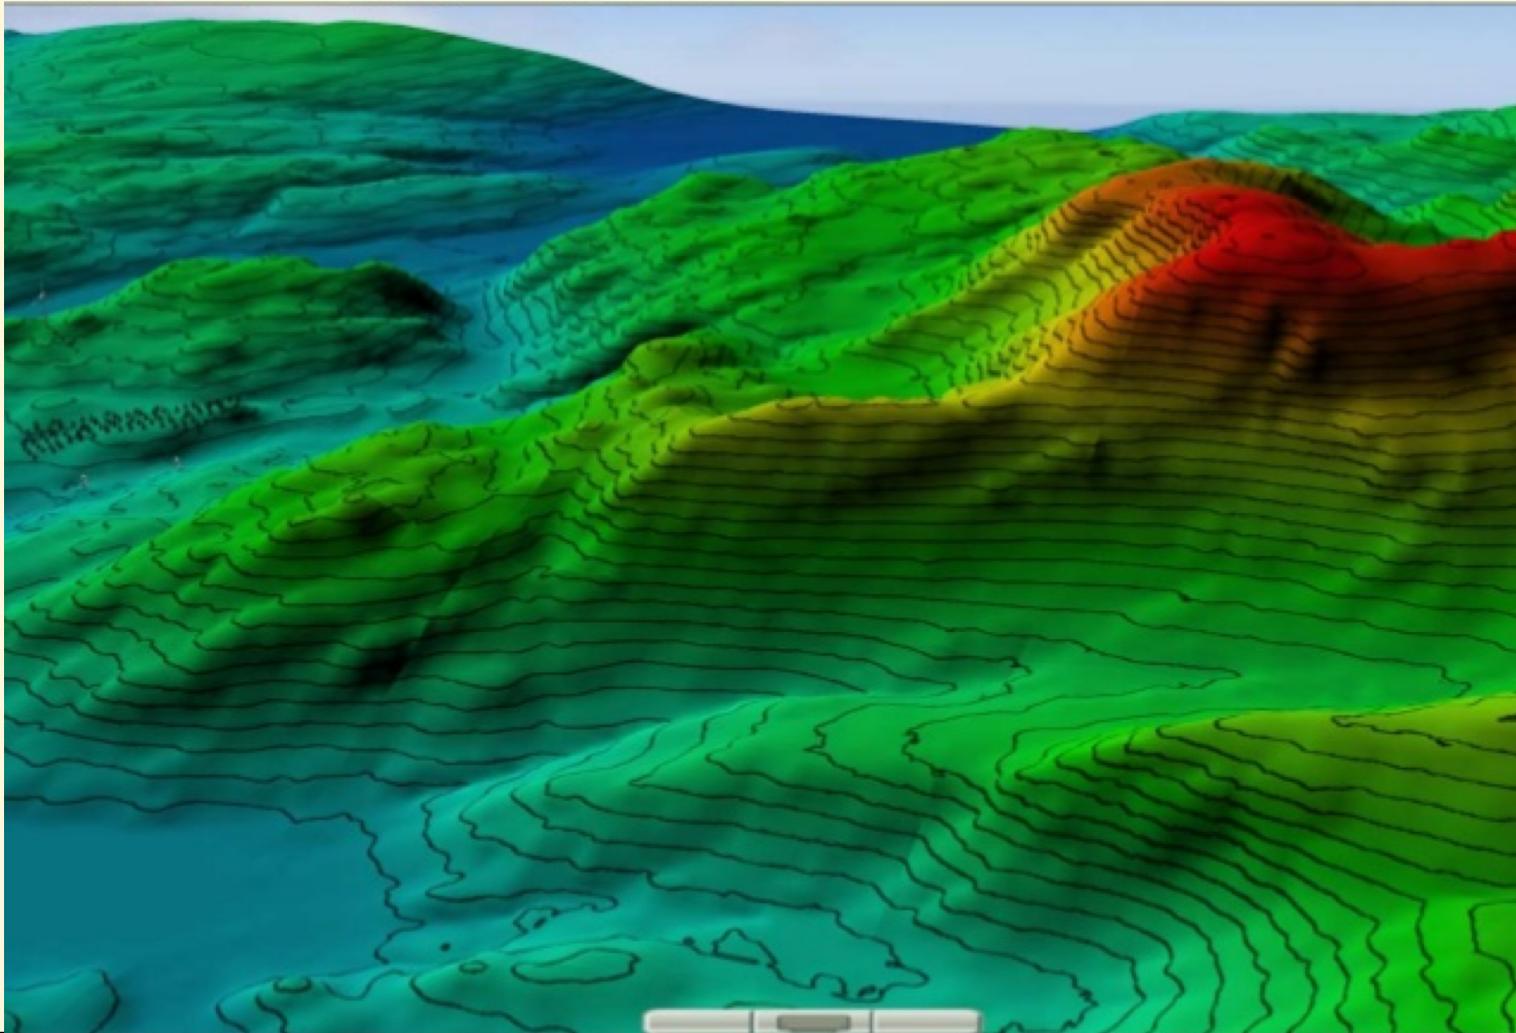
University of Colorado **Boulder**

How does this visualization help us?



University of Colorado **Boulder**

Terrain Concepts



University of Colorado **Boulder**

Terrain Concepts

Basic notions:

- Uphill / downhill
- Contour lines (curves of constant elevation)
- Steepness of slope
- Peaks/Valleys (local extrema)

More mathematical notions:

- Tangent Plane
- Normal vectors
- Curvature

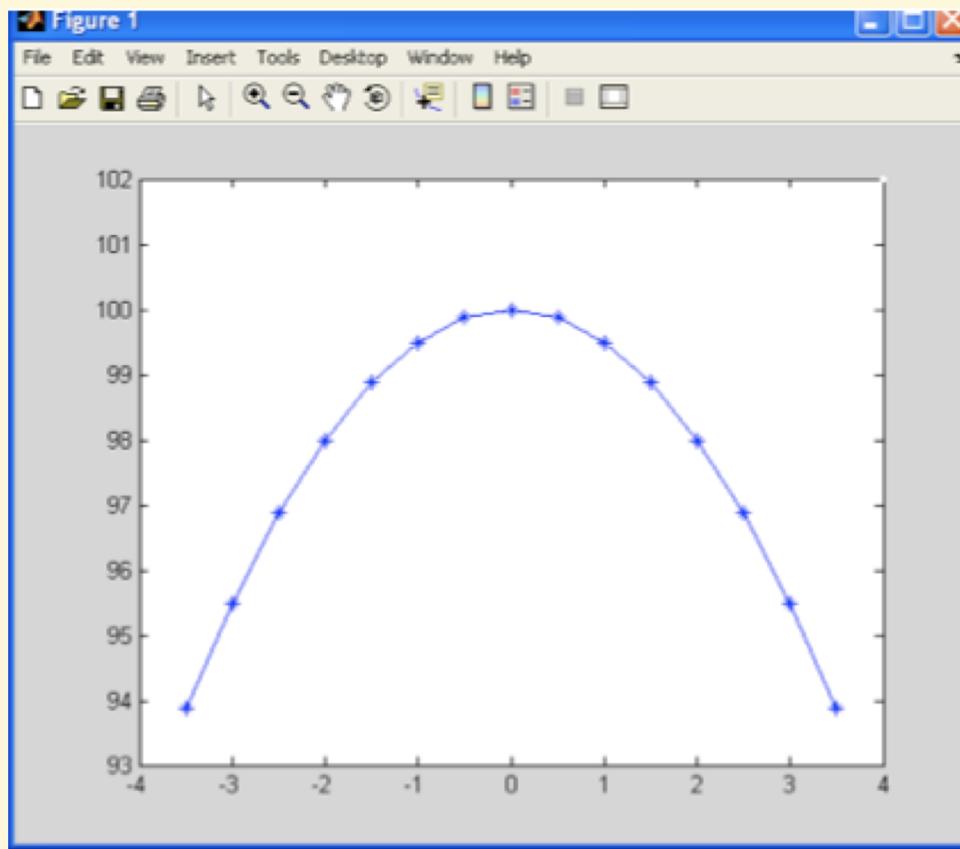
Gradient vectors (vectors of partial derivatives) will help us define/compute all of these.



University of Colorado **Boulder**

1 D Gradient

Consider function $f(x) = 100 - 0.5 * x^2$

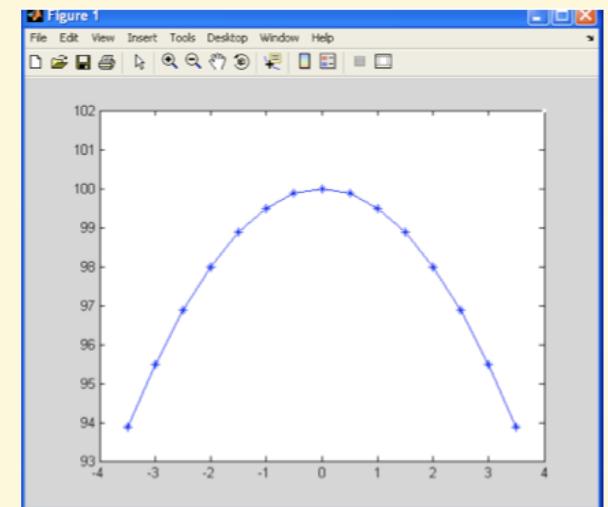
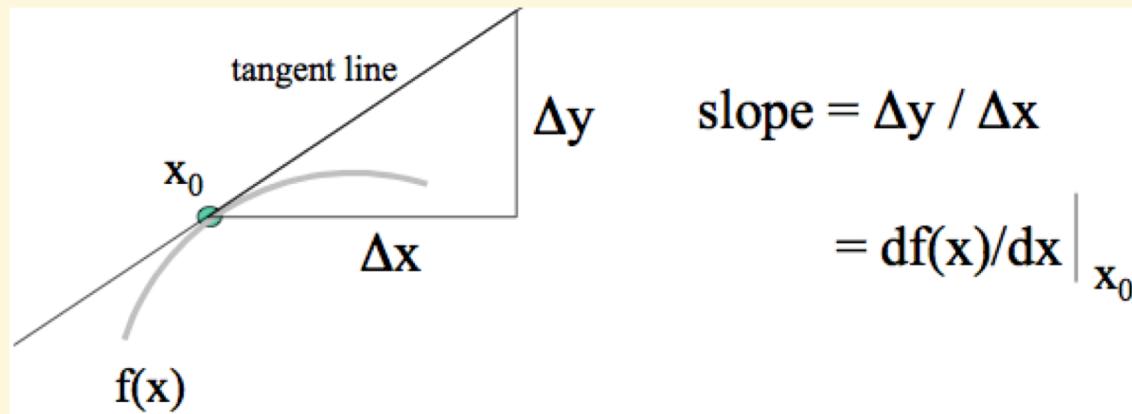


1 D Gradient

Consider function $f(x) = 100 - 0.5 * x^2$

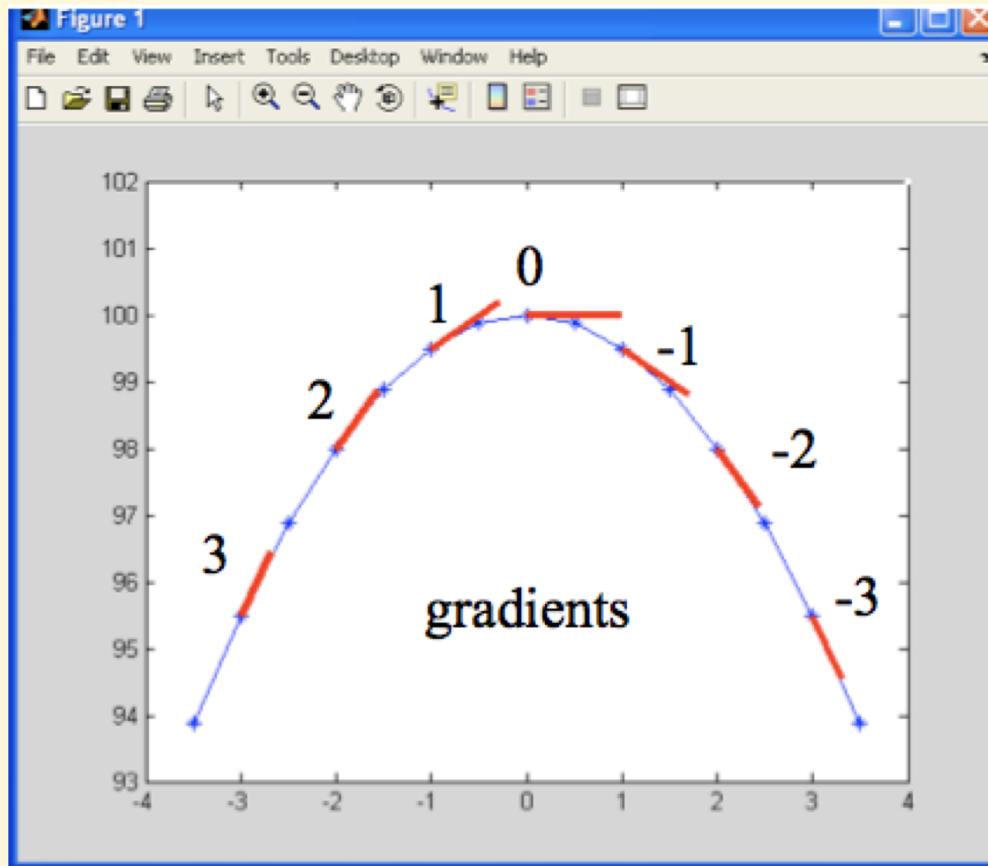
Gradient is $df(x)/dx = -2 * 0.5 * x = -x$

Geometric interpretation:
gradient at x_0 is slope of tangent line to curve at point x_0



1 D Gradient

$$f(x) = 100 - 0.5 * x^2 \quad df(x)/dx = -x$$

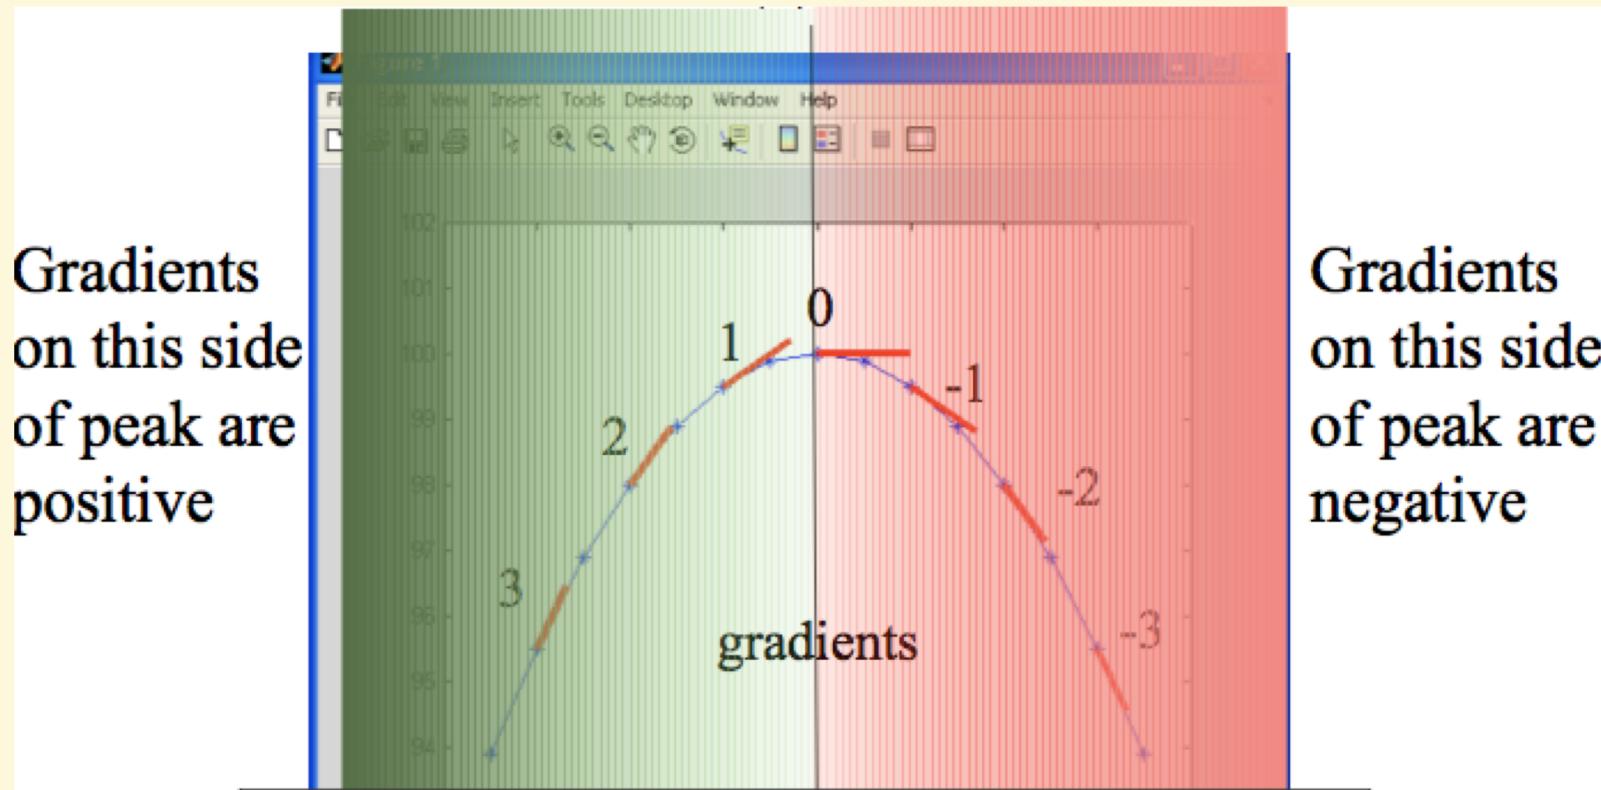


University of Colorado **Boulder**

1 D Gradient

$$f(x) = 100 - 0.5 * x^2$$

$$df(x)/dx = -x$$



Note: Sign of gradient at point tells you what direction to go to travel “uphill”



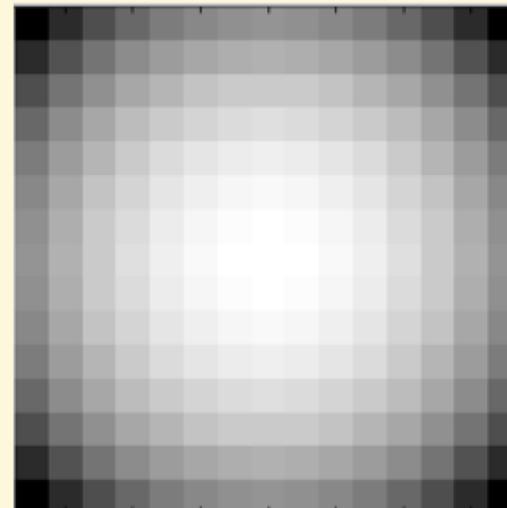
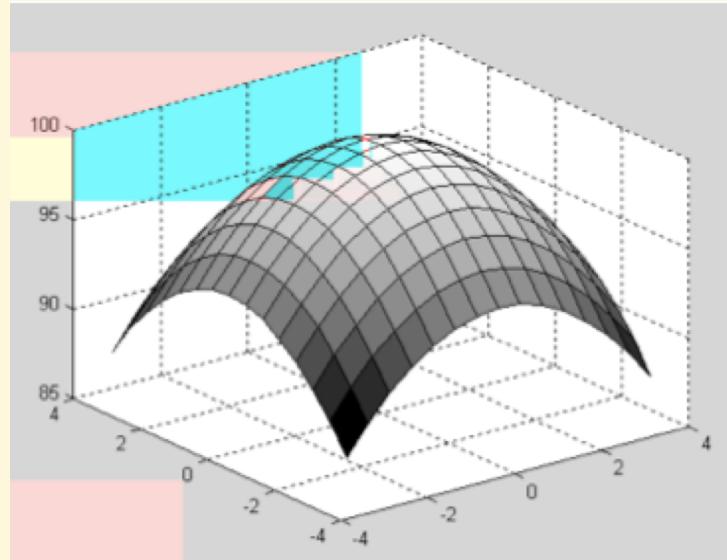
University of Colorado **Boulder**

2 D Gradient

$$f(x,y) = 100 - 0.5 * x^2 - 0.5 * y^2$$

$$\frac{df(x,y)}{dx} = -x \quad \frac{df(x,y)}{dy} = -y$$

$$\text{Gradient} = [\frac{df(x,y)}{dx}, \frac{df(x,y)}{dy}] = [-x, -y]$$



Gradient is vector of partial derivs wrt x and y axes

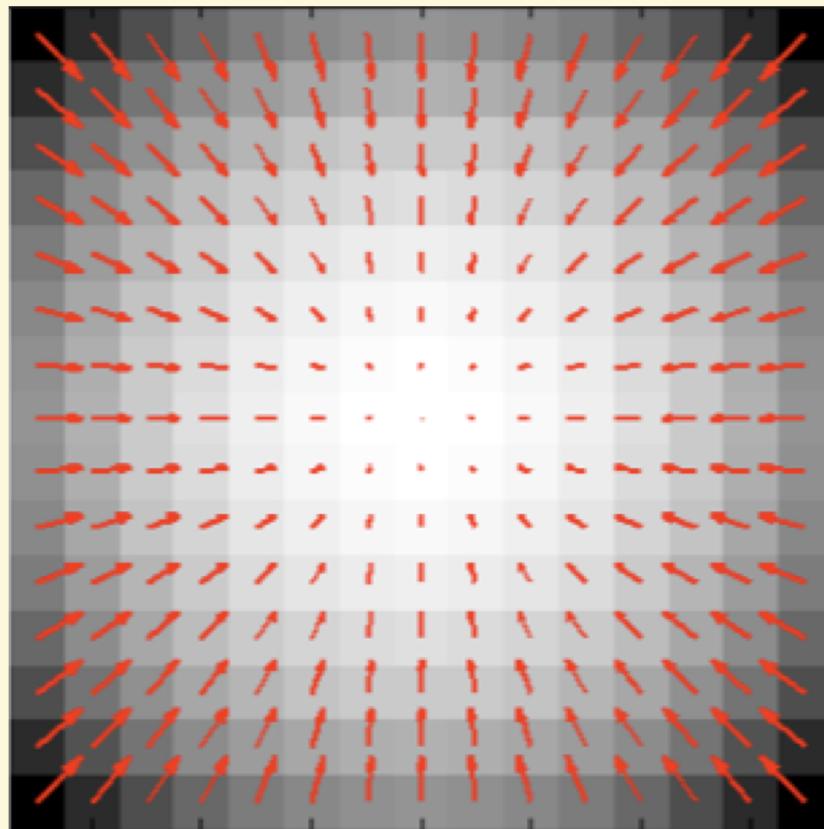


University of Colorado **Boulder**

2 D Gradient

$$f(x,y) = 100 - 0.5 * x^2 - 0.5 * y^2$$

$$\text{Gradient} = [df(x,y)/dx, df(x,y)/dy] = [-x, -y]$$



Plotted as a vector field,
the gradient vector at each
pixel points “uphill”.

The gradient indicates the
direction of steepest ascent.

The gradient is 0 at the peak
(also at any flat spots, and local minima)

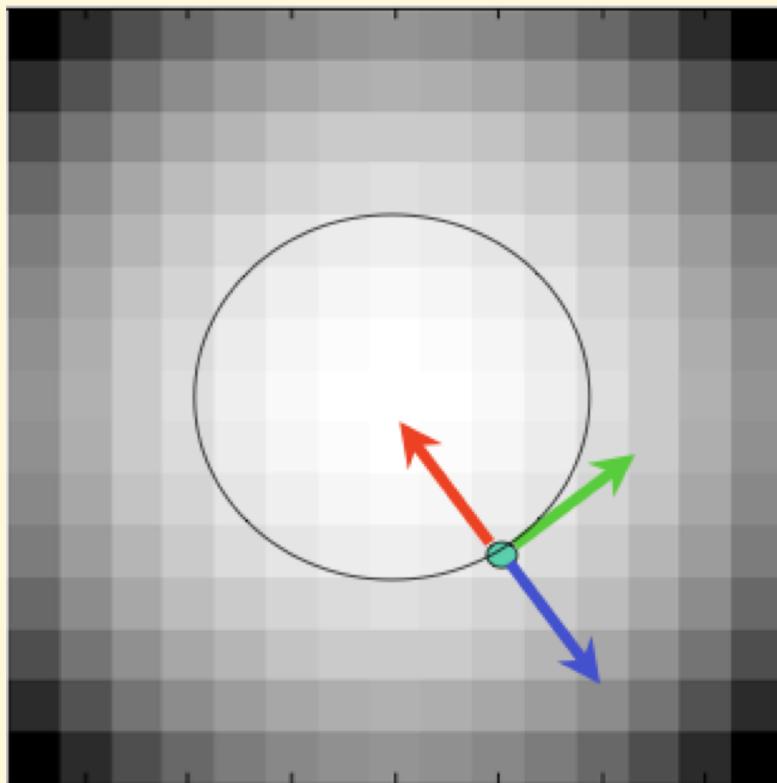


University of Colorado **Boulder**

2 D Gradient

$$f(x,y) = 100 - 0.5 * x^2 - 0.5 * y^2$$

$$\text{Gradient} = [df(x,y)/dx, df(x,y)/dy] = [-x, -y]$$



Let $\mathbf{g}=[g_x, g_y]$ be the gradient vector at point/pixel (x_0, y_0)

Vector \mathbf{g} points uphill
(direction of steepest ascent)

Vector $-\mathbf{g}$ points downhill
(direction of steepest descent)

Vector $[g_y, -g_x]$ is perpendicular,
and denotes direction of constant elevation. i.e. normal to contour
line passing through point (x_0, y_0)



Numerical Derivatives

Taylor's Expansion

$$f(x+h) = f(x) + f'(x)h + \frac{1}{2}f''(x)h^2 + O(h^3)$$

$$f'(x)h = f(x+h) - f(x) - \frac{1}{2}f''(x)h^2 + O(h^3)$$

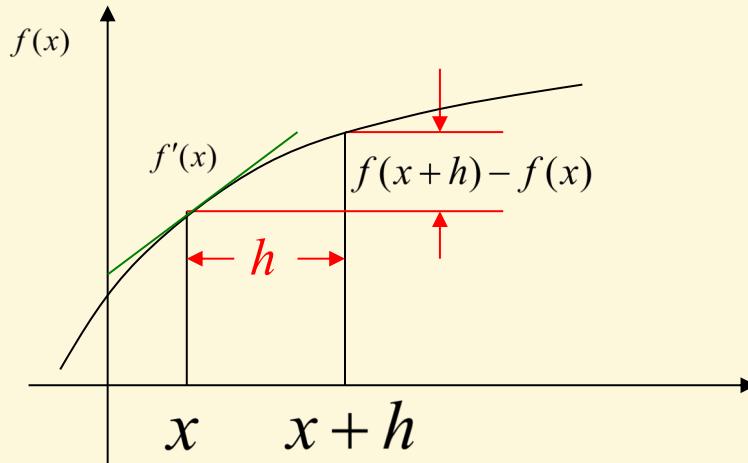
$$f'(x) = \frac{f(x+h) - f(x)}{h} - \frac{1}{2}f''(x)h + O(h^2)$$



Forward Difference Formula for $f'(x)$

$$f'(x) \approx \frac{f(x+h) - f(x)}{h} \quad \text{error} = O(h)$$

Geometrically



University of Colorado **Boulder**

Backward Difference Formula for $f'(x)$

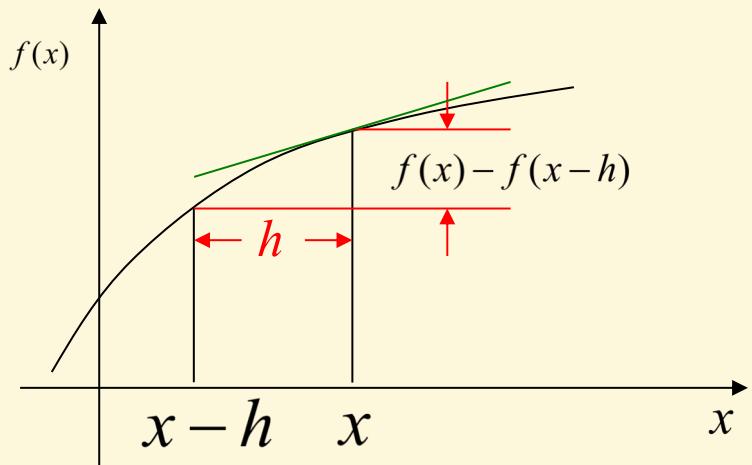
Similarly

$$f(x-h) = f(x) - f'(x)h + \frac{1}{2} f''(x)h^2 + O(h^3)$$

Geometrically

$$f'(x) \approx \frac{f(x) - f(x-h)}{h}$$

error = $O(h)$



University of Colorado **Boulder**

Central Difference Formula for $f'(x)$

$$f(x+h) = f(x) + f'(x)h + \frac{1}{2}f''(x)h^2 + \frac{1}{6}f'''(x)h^3 + \frac{1}{4!}f^{(4)}(x)h^4 + O(h^5)$$

$$-) \quad f(x-h) = f(x) - f'(x)h + \frac{1}{2}f''(x)h^2 - \frac{1}{6}f'''(x)h^3 + \frac{1}{4!}f^{(4)}(x)h^4 + O(h^5)$$

$$f(x+h) - f(x-h) = 2hf'(x) + \frac{1}{3}f'''(x)h^3 + O(h^5)$$

$$2hf'(x) = f(x+h) - f(x-h) - \frac{1}{3}f'''(x)h^3 + O(h^5)$$



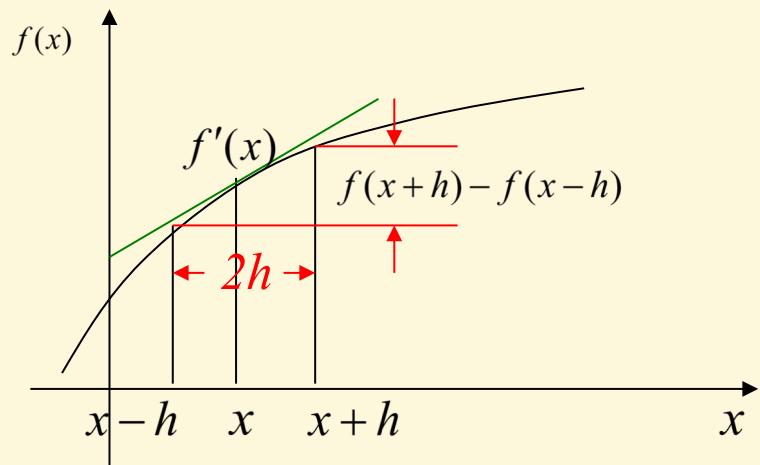
University of Colorado **Boulder**

Central Difference Formula for $f'(x)$

$$f'(x) \approx \frac{f(x+h) - f(x-h)}{2h}$$

error = $O(h^2)$

Geometrically



Example

$$f(x) = x^3 \quad (\therefore f'(x) = 3x^2, \quad f'(1) = 3)$$

Calculate $f''(1)$ using FD, BD, CD



Example (cont)

FD:

$$h=0.1 \quad f'(1) = \frac{f(1.1) - f(1)}{0.1} = 3.31 \quad \text{error} = 0.31$$

$$h=0.05 \quad f'(1) = \frac{f(1.05) - f(1)}{0.05} = 3.1525 \quad \text{error} = 0.1525$$

□ $\text{error} \propto h$

BD:

$$h=0.1 \quad f'(1) = \frac{f(1) - f(0.9)}{0.1} = 2.71 \quad \text{error} = 0.29$$

$$h=0.05 \quad f'(1) = \frac{f(1) - f(0.95)}{0.05} = 2.8453 \quad \text{error} = 0.1547$$

□ $\text{error} \propto h$

CD:

$$h=0.1 \quad f'(1) = \frac{f(1.1) - f(0.9)}{0.2} = 3.01 \quad \text{error} = 0.01$$

$$h=0.05 \quad f'(1) = \frac{f(1.05) - f(0.95)}{0.1} = 3.00250 \quad \text{error} = 0.00250$$

□ $\text{error} \propto h^2$



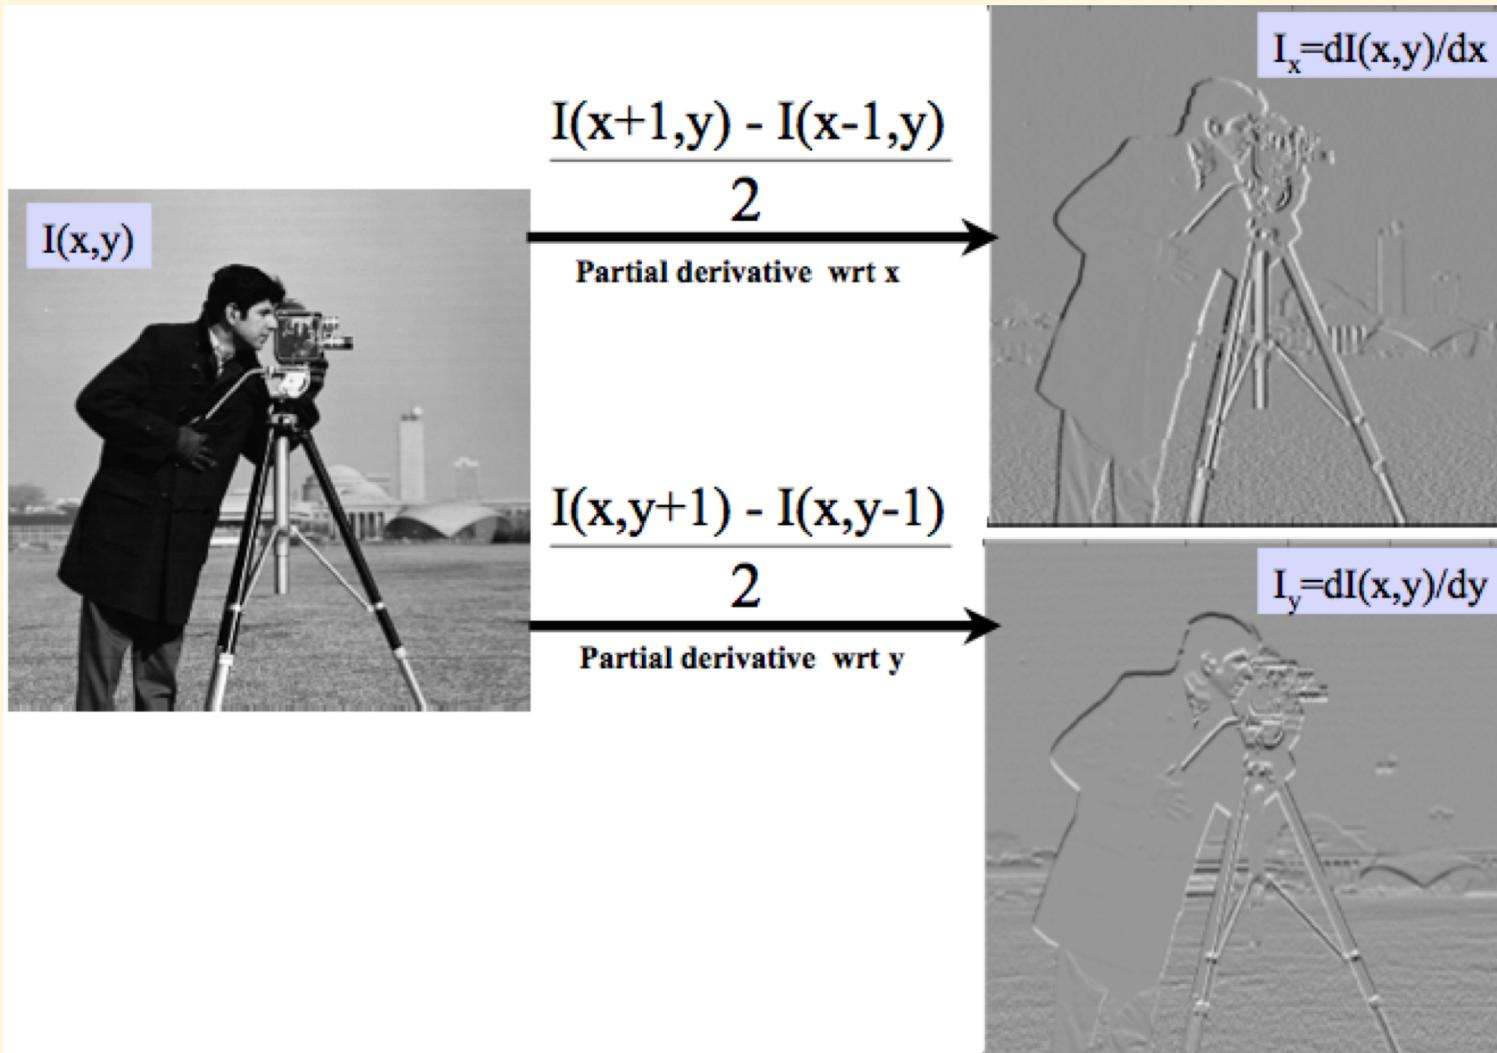
University of Colorado **Boulder**

Example (cont)

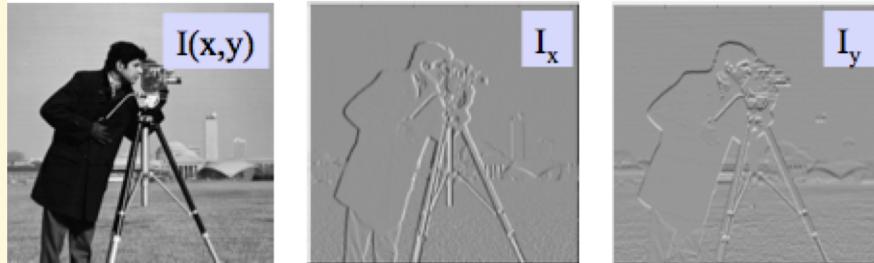
- Remarks:
 - FD, BD, CD each involves 2 function calls, 1 subtraction, and 1 division: same computation time
 - CD is the most accurate (hence, the most recommended method)
 - *However*, sometimes, CD cannot be applied



Example - Spatial Image Gradients



Functions of Gradients

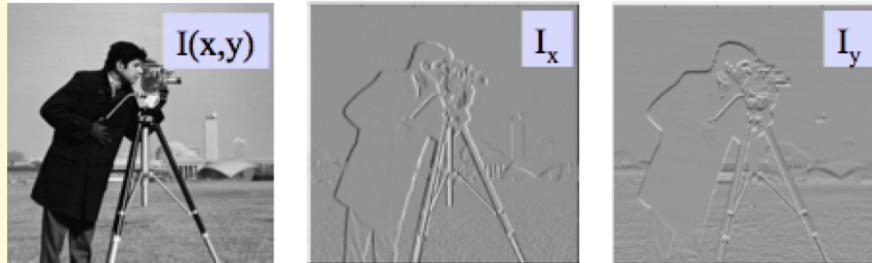


Magnitude of gradient
 $\sqrt{I_x.^2 + I_y.^2}$

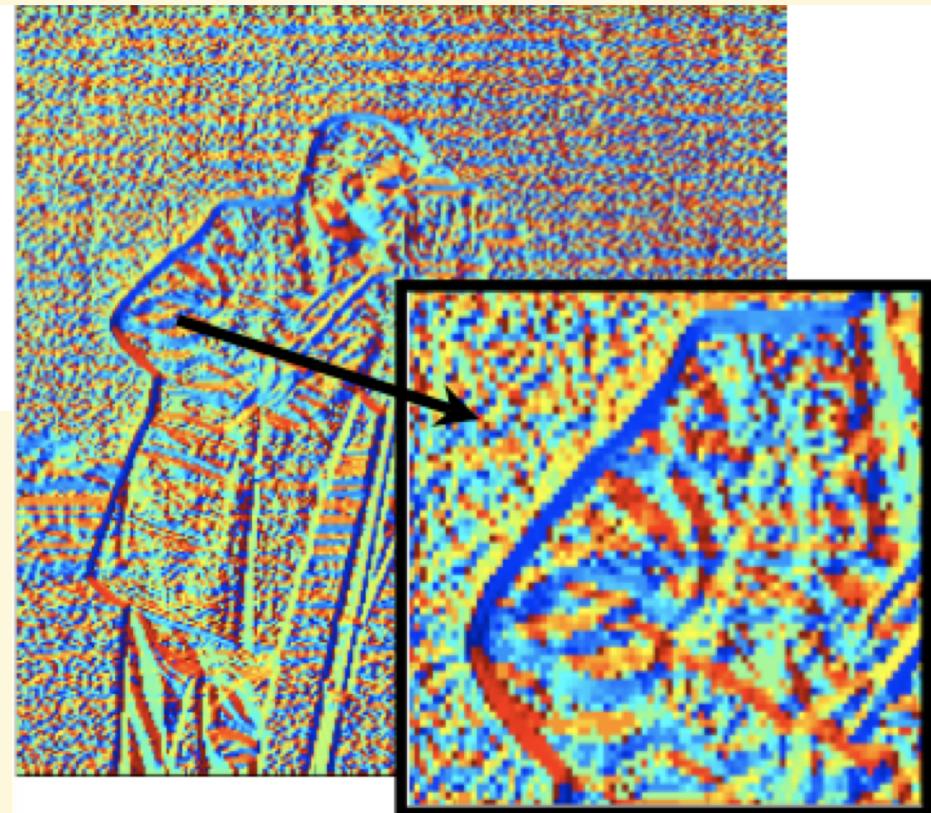


University of Colorado **Boulder**

Functions of Gradients

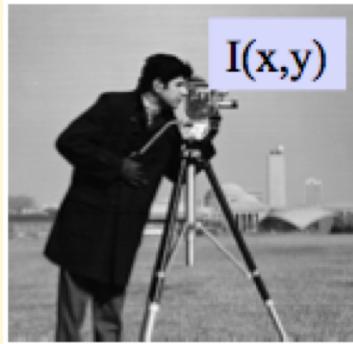


Angle of gradient
 $\text{atan2}(I_y, I_x)$



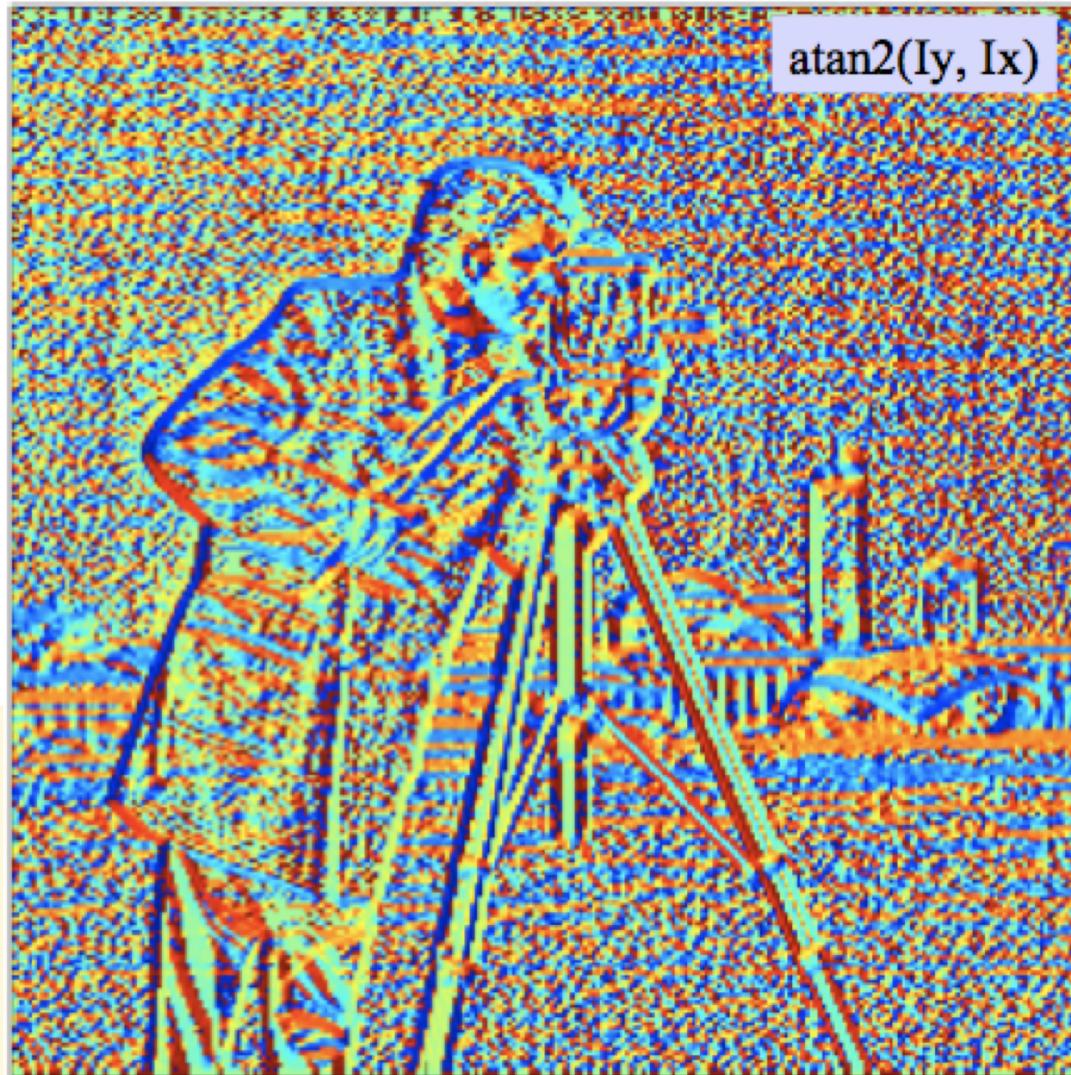
University of Colorado **Boulder**

Functions of Gradients



$I(x,y)$

What else do we observe in this image?



University of Colorado **Boulder**

Gradients as linear operators

Gradients are an example of linear operators, i.e. value at a pixel is computed as a linear combination of values of neighboring pixels.



University of Colorado **Boulder**

Image filtering

- Image filtering: compute function of local neighborhood at each position
- Really important!
 - Enhance images
 - Denoise, resize, increase contrast, etc.
 - Extract information from images
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching



University of Colorado **Boulder**

Example: box filter

$g[\cdot, \cdot]$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1



University of Colorado **Boulder**

Slide credit: David Lowe (UBC)

Image filtering

$$f[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[.,.]$$

$$g[.,.]$$

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$



University of Colorado **Boulder**

Credit: S. Seitz

Image filtering

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

$$h[\cdot, \cdot]$$

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

0	10									



University of Colorado **Boulder**

Credit: S. Seitz

Image filtering

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	90	0	0
0	0	0	90	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$



University of Colorado **Boulder**

$$h[\cdot, \cdot]$$

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Credit: S. Seitz

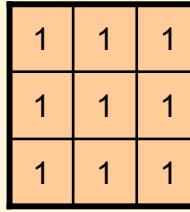
Image filtering

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

$$h[\cdot, \cdot]$$

$$g[\cdot, \cdot] \frac{1}{9}$$




University of Colorado **Boulder**

Credit: S. Seitz

Image filtering

 $f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0

 $h[\cdot, \cdot]$

1	1	1
1	1	1
1	1	1

$\frac{1}{9}$

			0	10	20	30	30				

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$



University of Colorado **Boulder**

Credit: S. Seitz

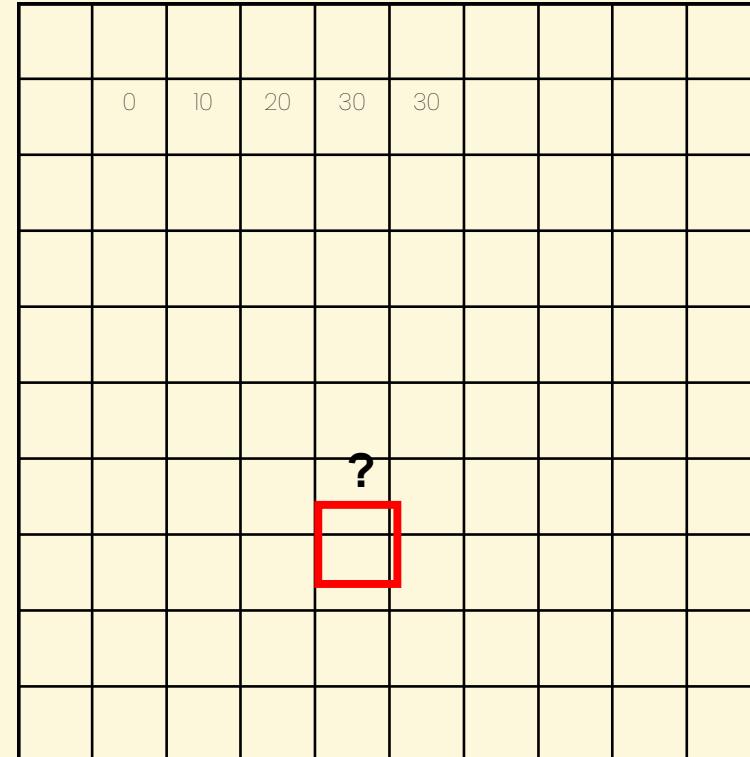
Image filtering

$f[\cdot, \cdot]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$g[\cdot, \cdot]$ $\frac{1}{9}$
 $h[\cdot, \cdot]$

1	1	1
1	1	1
1	1	1



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$



University of Colorado **Boulder**

Credit: S. Seitz

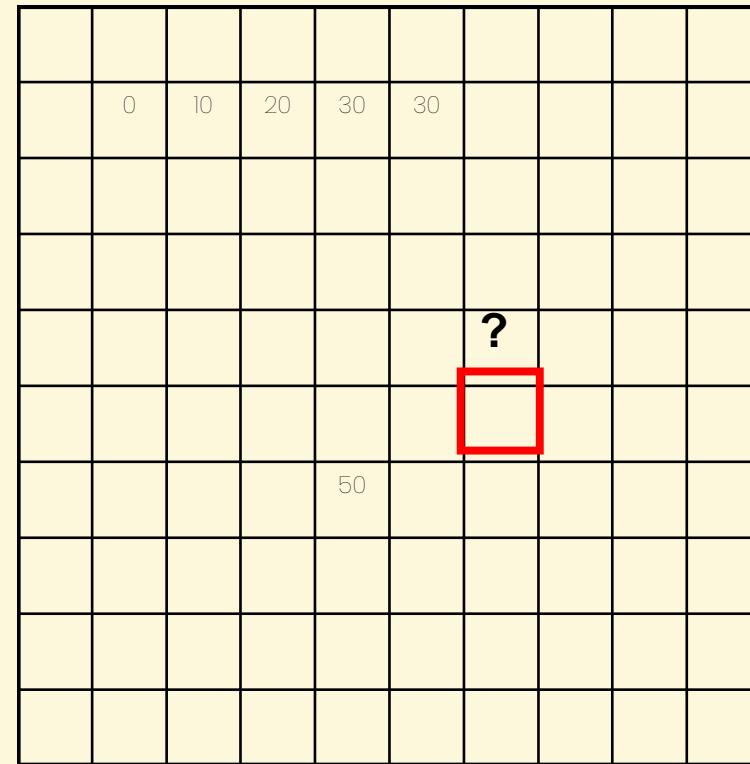
Image filtering

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h[\cdot, \cdot] \stackrel{\frac{1}{9}}{=} g[\cdot, \cdot]$$

1	1	1
1	1	1
1	1	1



$$h[m, n] = \sum_{k, l} g[k, l] f[m + k, n + l]$$



University of Colorado **Boulder**

Credit: S. Seitz

Image filtering

$$f[\cdot, \cdot]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$g[\cdot, \cdot] \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$$h[\cdot, \cdot]$$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
10	20	30	30	30	30	30	20	10		
10	10	10	0	0	0	0	0	0		

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$



University of Colorado **Boulder**

Credit: S. Seitz

Box Filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect (remove sharp features)

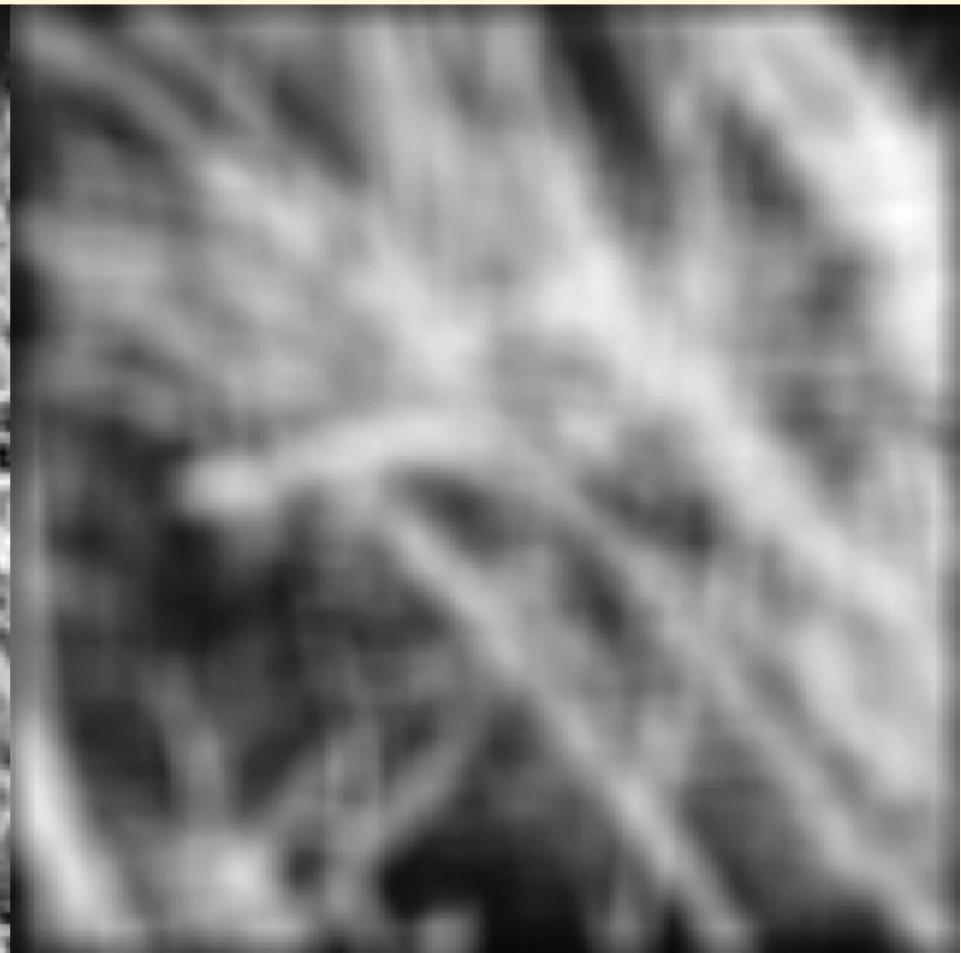
$g[\cdot, \cdot]$

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$


University of Colorado **Boulder**

Slide credit: David Lowe (UBC)

Smoothing with box filter



Practice with linear filters



0	0	0
0	1	0
0	0	0

?



University of Colorado **Boulder**

Source: D. Lowe

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)



University of Colorado **Boulder**

Source: D. Lowe

Practice with linear filters



0	0	0
0	0	1
0	0	0

?



University of Colorado **Boulder**

Source: D. Lowe

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel



University of Colorado **Boulder**

Source: D. Lowe

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)



University of Colorado **Boulder**

Source: D. Lowe

Practice with linear filters



Original

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

-

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



Sharpening filter

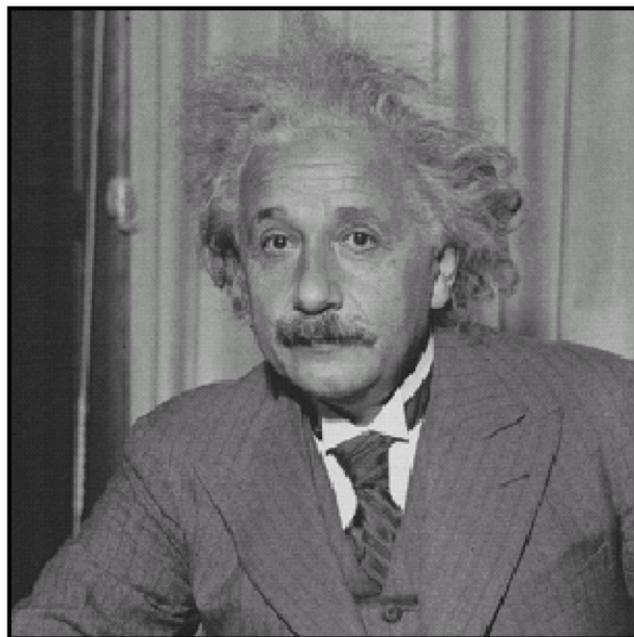
- Accentuates differences with local average



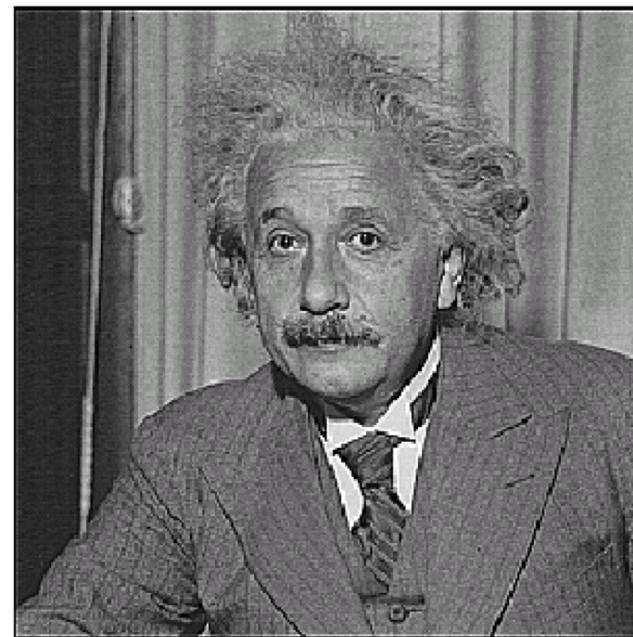
University of Colorado **Boulder**

Source: D. Lowe

Sharpening



before



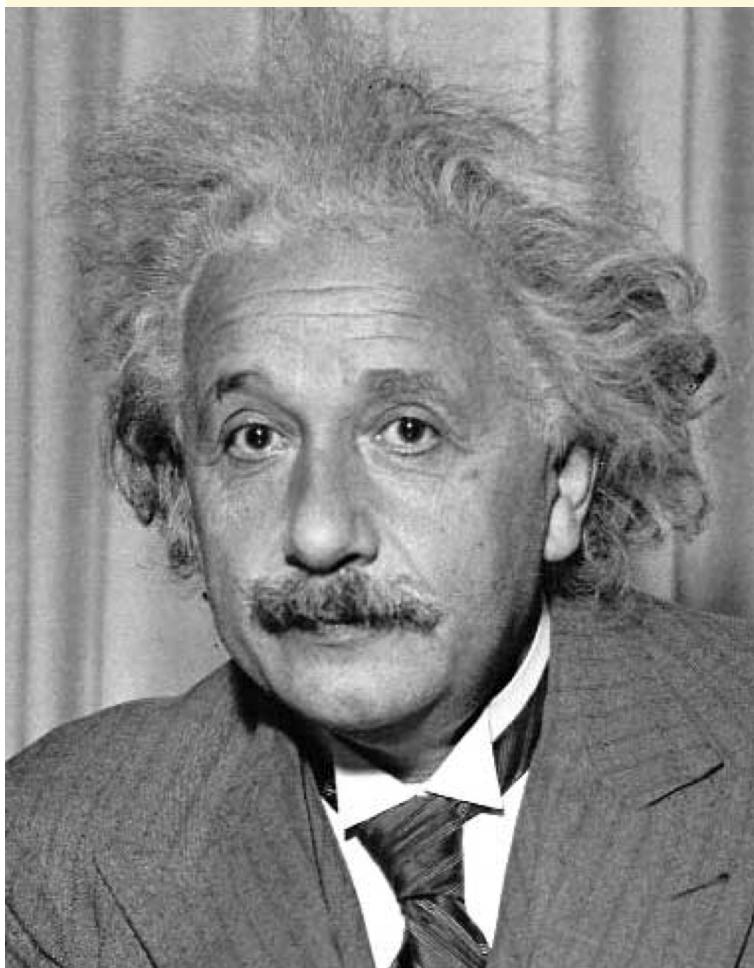
after



University of Colorado **Boulder**

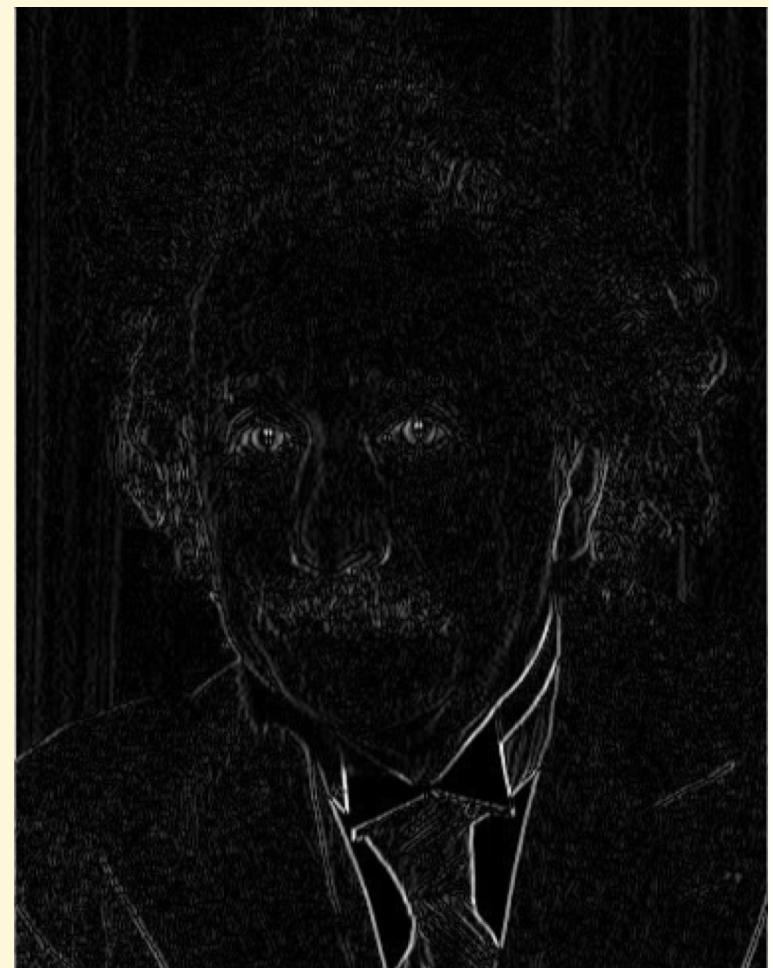
Source: D. Lowe

Other filters



1	0	-1
2	0	-2
1	0	-1

Sobel

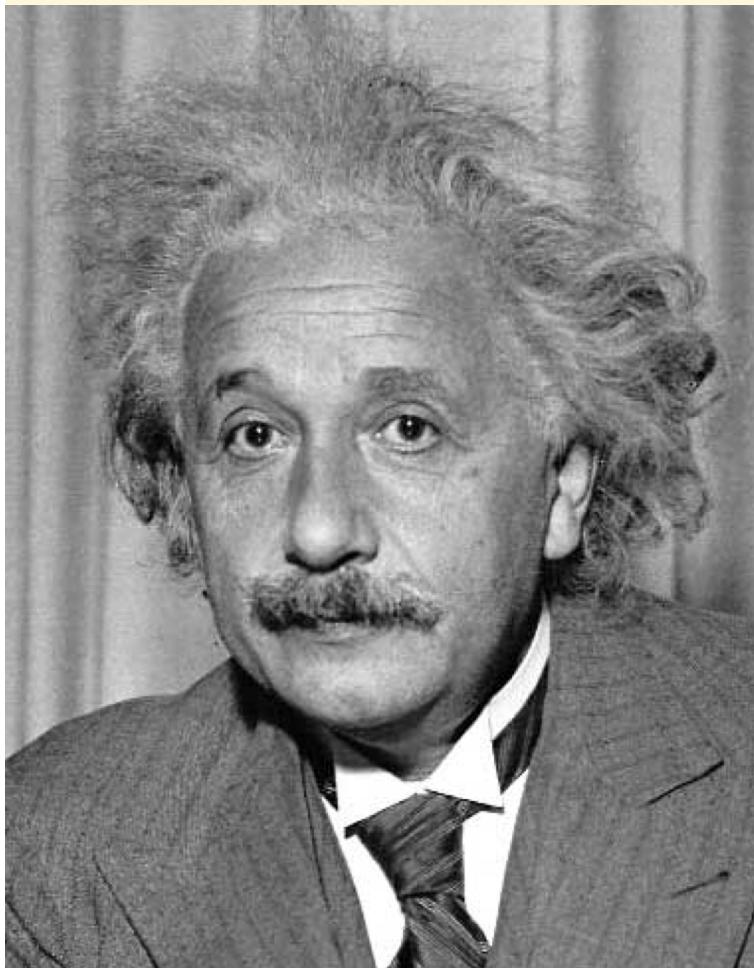


Vertical Edge
(absolute value)



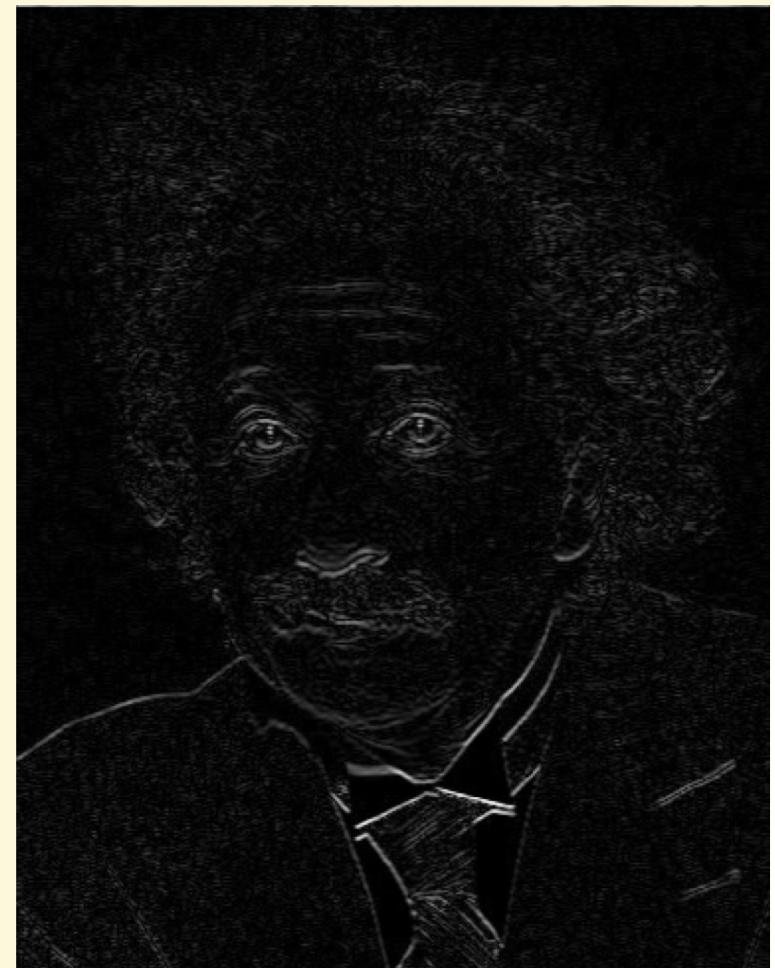
University of Colorado **Boulder**

Other filters



1	2	1
0	0	0
-1	-2	-1

Sobel



Horizontal Edge
(absolute value)



University of Colorado **Boulder**